

COSC 3750 — EXAM 1 STUDY GUIDE

What "System Programming" Mean (Big Picture)

- System programming is writing programs that talk directly to the OS:
 - Manage
 - Files
 - Processes
 - Permissions
 - Signals
 - Terminals
 - Network Connections
 - Use system calls (kernel interfaces) rather than just high-level libraries
 - Care about resources, errors, and edge cases
- Why Go Fits Well
 - Strong standard library
 - Files, Net, OS, IO
 - Explicit error handling → fewer failures
 - Goroutines & channels make concurrency safer than raw threads
- System Calls Vs. Library Calls
 - System Calls
 - Transition from user space → kernel
 - libc / Go Wrappers Exist For:
 - Portability, Convenience, Error Translation
 - Not every function you call is a system call

Go Basics for System Code (CH. 1-2)

- Toolchain
 - go build, go run, go test
- Modules manage dependencies
 - go.mod
- Code is organized in packages
- Idioms you MUST expect:
 - Always checks errors
 - Small, focused functions
 - Interfaces define behavior, not data

Files and File I/O (GO Ch. 3 + TLPI Ch. 4-5)

- Core Model
 - Files are accessed through file descriptors
 - Almost everything is treated as a file
 - Files, Pipes, Sockets, Terminals → treated like files
- Key Rules
 - open → read/write → close
 - reads and writes may be partial → loop until done
 - Always close resources
 - fd leaks = real bugs
- WHY Reads/Writes Happen
 - Pipes & Sockets
 - Nonblocking I/O
 - Kernel Buffering Limits
 - Signals Interrupting Syscalls
- Metadata
 - Files have:
 - size, permissions, owner, timestamps
 - A *filename* is just a directory entry, NOT the file itself

Permissions & Filesystems (Go Ch. 4 + TLPI Ch. 14-15)

- Permissions
 - Format
 - rwxrwxrwx (user / group / other)
 - umask
 - removes permissions from newly created files
 - chmod changes permissions | chown changes ownership
- Filesystems
 - A filesystem = data + metadata
 - Inode
 - Holds metadata
 - Permissions, owner, size
 - Filenames
 - Point to inodes → name ≠ file
 - Mounting
 - Merges filesystems into one directory tree

Directories & Links (TLPI Ch. 18)

- Directories map names → inode numbers
- Hard Link
 - Another name for the same inode
- Symbolic Link
 - A file that points to a pathname
- Deleting a file
 - unlink
 - Removes a name
 - Data is freed only when link count = 0 AND file not open

Go I/O Interfaces (Go Ch. 5)

- Stream Mental Model
 - Treat everything as a system of bytes
 - io.Reader / io.Writer
 - Unify files, sockets, buffers
- Composition
 - You can wrap streams
 - Buffering
 - Limiting
 - Tee (Copy to Multiple Outputs)
 - This enables powerful pipelines with minimal code

Concurrency (Go Ch. 6)

- Goroutines
 - Lightweight concurrent functions
- Channels
 - Communication & Synchronization
- Use concurrency when
 - You need parallel work
 - You need to wait on multiple events
- Still need coordination for
 - Ordering
 - Cancellation
 - Limited Resources
- Goroutines ≠ OS Threads
 - Go schedules goroutines onto threads
 - Much cheaper to create and manage

Signals (Go Ch. 7 + TLPI Ch. 20-22)

- What Signals Are
 - Asynchronous Notifications
 - Ctrl-C, Timers, Child Exit
 - Default Actions
 - Terminate, Stop, Ignore, Continue
- Signals Can Interrupt:
 - read, write, accept
- Programs MUST:
 - Retry OR handle EINTR
- Important Concepts
 - Generated Vs. Delivered
 - Pending Vs. Handling
 - Signal Masks Block Delivery
- Most functions are NOT async-signal-safe
- Handlers
 - Interrupt normal execution
 - MUST be minimal & async-signal-safe
 - In GO
 - Usually delivered through channels → clean shutdown paths

Processes & Execution (TLPI Ch. 24-35)

- fork / exec Model
 - fork()
 - Child is a copy
 - Same memory, same fds
 - exec()
 - Replace program image, Same PID
- Termination
 - Processes exit with a status
 - Parent must wait() or waitpid() to avoid zombies
- Why wait() Matters
 - Without wait()
 - Zombie processes accumulate
 - With wait()
 - Reaps Child
 - Retrieves Exit Status
- Zombies vs. Orphans
 - Zombie: exited, not waited on
 - Orphan: parent exited; adopted by init/systemd

Process Groups, Job Control (TLPI Ch. 34)

- Processes belong to process groups
- Groups belong to sessions
- Terminal send signals (Ctrl-C) to the foreground process group
- Job control exists so shells can manage foreground/background jobs

Networking Basics (Go Ch. 8)

- Client/Server Model
- Common Pattern:
 - Listen → Accept → Read/Write → Close
- Go's `net` package abstracts sockets cleanly

Daemons & Logging (TPLI Ch. 37)

- Daemons:
 - Run in background
 - No controlling terminal
 - Clean environment
 - cwd, umask, file descriptors
- Logging goes to log files or syslog, NOT stdout

Terminals (TLPI Ch. 62)

- Terminal Driver Model
 - Input/Output Queues & Line Discipline
 - Terminal interprets special characters
- Canonical Vs. Noncanonical
 - Canonical:
 - line-at-a-time
 - Enter REQUIRED
 - Noncanonical:
 - character-at-a-time
 - Used by editors, games
- Special Characters
 - INTR → Ctrl-C → SIGINT
 - EOF → Ctrl-D
 - SUSP → Ctrl-Z
- Termios
 - `tcgetattr / tcsetattr` control terminal behavior
 - Tools like `stty` modify these flags

Pseudoterminals (TLPI Ch. 64)

- **pty slave**
 - Looks like a real terminal
- **pty master**
 - Controlled by another program
- **Used by**
 - Shells, ssh, script, terminal emulators
- **WHY SSH, Shells, Scripts Need PTYs**
 - Programs behave differently when stdin is a terminal
 - Line editing, echoing, signals require terminal semantics
 - ptys fake a real terminal
- **Typical Pattern**
 1. Open pty master
 2. fork
 3. Child: New session → open slave → dup to stdin/out/err → exec
 4. Parent: relay data between user and child

Pipes (TLPI Ch. 44)

- Unidirectional byte streams
- Used heavily by shells
- EOF happens when write end is closed
 - Occur only when ALL write ends are closed
- Deadlocks happen when both sides wait forever
- Deadlocks Happens When:
 - One process waits to read
 - Other waits to write
 - Neither closes unused pipe ends
- ALWAYS close unused ends after fork to avoid deadlocks
- Classic Pattern
 - pipe → fork → dup → exec

Exam Habits to Remember

- Assume partial reads/writes
- Always know what must be closed
- Files, Pipes, Sockets, Terminals ALL share the same I/O model
- Errors are normal → handle them
- Think in terms of resources and lifetimes

File Descriptor Flags Vs File Status Flags

- **File Descriptor Flags**
 - Per descriptor
 - Ex: close-on-exec
- **File Status Flags**
 - Per open file
 - EX: Append, Nonblocking
- **They affect behavior differently & live in different places**

Mental Model That Solves Most Questions – Ask Yourself

- **What resource is created?**
- **Who owns it?**
- **Who must close it?**
- **What happens if it isn't closed?**